

A proposal of new concurrency model for Ruby 3

Koichi Sasada
ko1@heroku.com



People love “Concurrency”



dRuby in the last century.

Masatoshi SEKI @m_seki

JA



Improved scalability by relaxing the GVL

Charlie Gracie @crgracie

EN



ErRuby: Ruby on Erlang/OTP

Lin Yu Hsiang @johnlinvc

EN



concurrent-ruby and how it is making Rails concurrent

Vipul A M @vipulnsward

EN



How to create multiprocess server on Windows with Ruby

Ritta Narita @narittan

EN



Ruby Concurrency compared

Anil Wadghule @anildigital

EN

Concurrent RubyKaigi

*(at least, there are two
parallel sessions)*

Why people love (to discuss)
“Concurrency”?

- Performance by “Parallel” execution to utilize multiple-cores

About this presentation

- Show “Why difficult multi-threads programs”
- Propose new concurrent and parallel mechanism idea named **“Guild”**
 - For Ruby 3

Koichi Sasada

- A programmer living in Tokyo, Japan
- Ruby core committer since 2007
 - YARV, Fiber, ... (Ruby 1.9)
 - RGenGC, RincGC (Ruby 2...)



Koichi is an Employee



heroku

Difficulty of Multi-threads programming

Programming language evolution

- Trade-off: Performance v.s. Safety/Easily
 - Performance: making faster programs
 - Safety: making bug-free programs
 - Easily: making programs happily

Two example C language

- String manipulation with pointers
- Memory management without GC

String manipulation with pointers

- C: Using raw pointers to manipulate pointers
 - Good: all-purpose and fast
 - Bad: easy to introduce bugs
 - Generate strange behavior
 - Cause abnormal termination
- Ruby: Wrap with String class
 - Good: Easy to use
 - Bad: slower than C in some cases

Object management without GC

- C: Free memory objects manually
 - Good: full control (target, timing and so on)
 - Bad: easy to introduce bugs (double-free/memory-leak, ...)
- Ruby: Automatic collection with GC
 - Good: nothing to care about object collection
 - Bad: introduce some overhead

Catch the two examples

- Ruby chose “safety/easily” approach
 - I believe nobody want to use pointers to manipulate String objects and want to manage objects manually anymore
 - Because ...
 - Nowadays computers is enough faster.
 - Implementation techniques overcome performance penalties (like generational GC and so on)

Current concurrent primitive Multi-threads programming

- Multi-threads programming is difficult

- **Introduce data race, race condition**

- Introduce deadlock, livelock

- Difficulty on debugging because of nondeterministic behavior

- difficult to reproduce same problem

**Difficult to make
correct (bug-free)
programs**

- Difficult to tune performance

**Difficult to make
fast programs**

Data race and race condition

- Bank amount transfer example
 - Quoted from Race Condition vs. Data Race
<http://blog.regehr.org/archives/490>

```
def transfer1 (amount, account_from, account_to)
  if (account_from.balance < amount) return NOPE
  account_to.balance += amount
  account_from.balance -= amount
  return YEP
end
```

Data race

- “`account_to.balance += amount`” has **Data-race**
 - Assume two threads (T1 and T2) invoke this methods with same bank accounts

interleave two threads (T1: amount = 100, T2: amount = 200)

T1: `t1 = account_to.balance` # `t1 = 10`

T2: `t2 = account_to.balance` # `t2 = 10`

T1: `account_to.balance = t1 + 100` \Rightarrow 110

T2: `account_to.balance = t2 + 200` \Rightarrow 210 (**expected: 310**)

Race condition

- To avoid data-race with the lock
- But there is another problem yet

```
# Lock with "Thread.exclusive"
def transfer2 (amount, account_from, account_to)
  if (account_from.balance < amount) return NOPE
  Thread.exclusive{ account_to.balance += amount }
  Thread.exclusive{ account_from.balance -= amount }
  return YEP
end
```

Race condition

- To avoid data-race with the lock
- But there is another problem yet

```
# T1 amount = 100, T2 amount = 200, account_from.balance = 250
T1: if (account_from.balance (== 250) < 100) return NOPE # OK, go through
T2: if (account_from.balance (== 250) < 200) return NOPE
T2: Thread.exclusive{ account_to.balance += 200 }
T2: Thread.exclusive{ account_from.balance -= 200 } #=> 250-200 => 50
T1: Thread.exclusive{ account_to.balance += 100 }
T1: Thread.exclusive{ account_from.balance -= 100 } #=> 50 - 100 => negative number!!
```

Final solution

- Lock whole of method

```
def transfer1 (amount, account_from, account_to)
  Thread.exclusive{
    if (account_from.balance < amount) return NOPE
    account_to.balance += amount
    account_from.balance -= amount
    return YEP
  }
end
```

Another example

Multi-thread quiz

- What happen on this program?

```
ary = [1, 2, 3]
t1 = Thread.new{
  ary.concat [4, 5, 6]
}
t2 = Thread.new{
  p ary # what's happen?
}.join
```

- (1) [1, 2, 3]
- (2) [1, 2, 3, 4, 5, 6]
- (3) (1) or (2)

Another example

Multi-thread quiz

- Answer: (4) depends on an interpreter

```
ary = [1, 2, 3]
t1 = Thread.new{
  ary.concat [4, 5, 6]
}
t2 = Thread.new{
  p ary # what's happen?
}.join
```

On MRI, (3) is correct

It will shows

[1, 2, 3] or

[1, 2, 3, 4, 5, 6]

(depends on thread
switching timing)

Another example

Multi-thread quiz

- Answer: (4) depends on interpreter

```
ary = [1, 2, 3]
t1 = Thread.new{
  ary.concat [4, 5, 6]
}
t2 = Thread.new{
  p ary # what's happen?
}.join
```

On JRuby:

It can cause Java
exception because
“Array#concat” is not
thread safe

On JRuby ...

```
# similar program
h = Hash.new(0)
NA = 1_000
10_000.times{
  ary = []
  (1..10).each{
    Thread.new{
      NA.times{|i|
        ary.concat [i]
      }
    }
  }
  t2 = Thread.new{
    s = ary.dup
  }.join
}
```



Unhandled Java exception: java.lang.NullPointerException

```
java.lang.NullPointerException: null
  rbInspect at org/jruby/RubyBasicObject.java:1105
  inspect at org/jruby/RubyObject.java:516
  inspectAry at org/jruby/RubyArray.java:1469
  inspect at org/jruby/RubyArray.java:1497
  cacheAndCall at org/jruby/runtime/callsite/CachingCallSite.java:293
  call at org/jruby/runtime/callsite/CachingCallSite.java:131
  block in t.rb at t.rb:17
  yieldDirect at org/jruby/runtime/CompiledIRBlockBody.java:156
  yieldSpecific at org/jruby/runtime/IRBlockBody.java:73
  yieldSpecific at org/jruby/runtime/Block.java:136
    times at org/jruby/RubyFixnum.java:291
  cacheAndCall at org/jruby/runtime/callsite/CachingCallSite.java:303
  callBlock at org/jruby/runtime/callsite/CachingCallSite.java:141
  call at org/jruby/runtime/callsite/CachingCallSite.java:145
  <top> at t.rb:3
  invokeWithArguments at java/lang/invoke/MethodHandle.java:599
    load at org/jruby/ir/Compiler.java:111
  runScript at org/jruby/Ruby.java:833
  runScript at org/jruby/Ruby.java:825
  runNormally at org/jruby/Ruby.java:760
  runFromMain at org/jruby/Ruby.java:579
  doRunFromMain at org/jruby/Main.java:425
  internalRun at org/jruby/Main.java:313
    run at org/jruby/Main.java:242
  main at org/jruby/Main.java:204
```

Difficulty of multi-threads programs

- We need to synchronize all sharing mutable objects correctly (or completely)
 - We need to know which methods are thread-safe and which methods are not.
 - Easy to track all of shared mutable objects on small program (like previous examples)
 - Difficult to track on big programs, especially on programs using gems
- We need to check all of source codes, or believe library documents (and we need to keep maintain documents correctly)

Difficulty of multi-threads programs (cont.)

- It requires “completeness”
 - If we forget to synchronize for only one mutable shared object, it will cause unexpected behavior.
- For debugging, it is difficult to find out the bugs
 - Backtrace may not work well because the problem may be placed on another line.
 - Bugs don't appear frequently with small data (test-data)
 - Difficult to reproduce issue because of nondeterministic behavior (depends on timing, etc).

FYI:

Why MRI Array#concat is thread-safe?

- MRI uses GVL (Giant/Global VM Lock) to control thread switching timing and C methods (such as Array#concat) are working atomically.
- GVL prohibits parallel thread execution (BAD), however it avoids several severe issues (GOOD).

Thread programming:

Performance tuning issue

```
a1 = []; a2 = []  
NA = 10_000_000  
t1 = Thread.new{  
  NA.times{|i| a1 << i }  
}.join  
t2 = Thread.new{  
  NA.times{|i| a2 << i }  
}.join
```

Serial program:

real 0m8.568s

user 0m37.816s

sys 0m5.530s

on JRuby

Thread programming: Performance tuning issue

```
a1 = []; a2 = []  
NA = 10_000_000  
t1 = Thread.new{  
  NA.times{|i| a1 << i }  
}  
t2 = Thread.new{  
  NA.times{|i| a2 << i }  
}  
t1.join; t2.join
```

Parallel program
(2 threads):

<u>real</u>	0m6.411s
user	0m20.527s
sys	0m7.798s

Thread programming:

Performance tuning issue

```
a1 = []; a2 = []
NA = 10_000_000
m1, m2 = Mutex.new, Mutex.new
t1 = Thread.new{
  NA.times{|i| m1.synchronize{ a1 << i }}
}
t2 = Thread.new{
  NA.times{|i| m2.synchronize{ a2 << i }}
}
t1.join; t2.join
```

Parallel program with
a useless lock 1
(2 threads):

<u>real</u>	0m10.264s
user	0m38.370s
sys	0m4.406s

Thread programming: Performance tuning issue

```
a1 = []; a2 = []  
NA = 10_000_000  
m = Mutex.new  
t1 = Thread.new{  
  NA.times{|i| m.synchronize{ a1 << i }}  
}  
t2 = Thread.new{  
  NA.times{|i| m.synchronize{ a2 << i }}  
}  
t1.join; t2.join
```

Parallel program with
a useless lock 2
(2 threads):

real 0m15.163s
user 0m45.317s
sys 0m9.658s

Performance tuning issue

	Execution time
Serial program	<u>8.568s</u>
Parallel program	<u>6.411s</u>
Parallel program with a useless lock 1	<u>10.264s</u>
Parallel program with a useless lock 2	<u>15.163s</u>

Thread programming: Performance tuning issue

We need to use just correct number locks

Not enough → unexpected behavior

Too much → performance penalty

FYI: synchronization mechanism

- Many synchronization mechanisms...
 - Mutual exclusion (Mutex), monitor, critical section
 - Transactional memory (optimistic lock)
 - Atomic instructions
 - Synchronized Queue
 - ...
 - Research on many lightweight lock algorithms
- They assume we can use them correctly

Overcome thread difficulty

Key idea

Problem:

Easy to share mutable objects

Idea:

Do not allow to share mutable objects
between concurrent entities (like
threads), **without any restriction**

Study from other languages

- Shell script with pipes, Racket (Place)
 - Copy mutable data between processes w/ pipes
- Erlang/Elixir
 - Do not allow mutable data
- Clojure
 - Basically do not allow mutable data
 - Special data structure to share mutable objects
 - Note that it can share mutable objects on Java layer

NOTE: we do not list approaches using “type system”

Don't you know
Elixir language?

Programming Elixir 1.2 by Dave Thomas

邦訳: プログラミング Elixir
笹田耕一・鳥井雪共訳 2016/08/19

You can buy it TODAY!!
サイン会は明日13時らしいです



Summary of approaches

- Communication with copied data (shell scripts)
 - Good: we don't need locks
 - Bad: copy everything is slow
- Prohibit mutable objects
 - Good: we don't need locks
 - Bad: Ruby utilizes many “write” operations. Unacceptable.
- Provide special data structure to share mutable objects
 - Good: we don't need locks (who don't use such special data structures)
 - Bad: Difficult to use special data structures.

Background was finished

Our goal for Ruby 3

- **We need to keep compatibility** with Ruby 2.
- We shouldn't consider about locks any more.
- We can share objects with copy, but copy operation should be fast.
- We should share objects if we can.
- We can provide special objects to share mutable objects like Clojure if we really need speed.



“Guild”

New concurrency model for Ruby 3

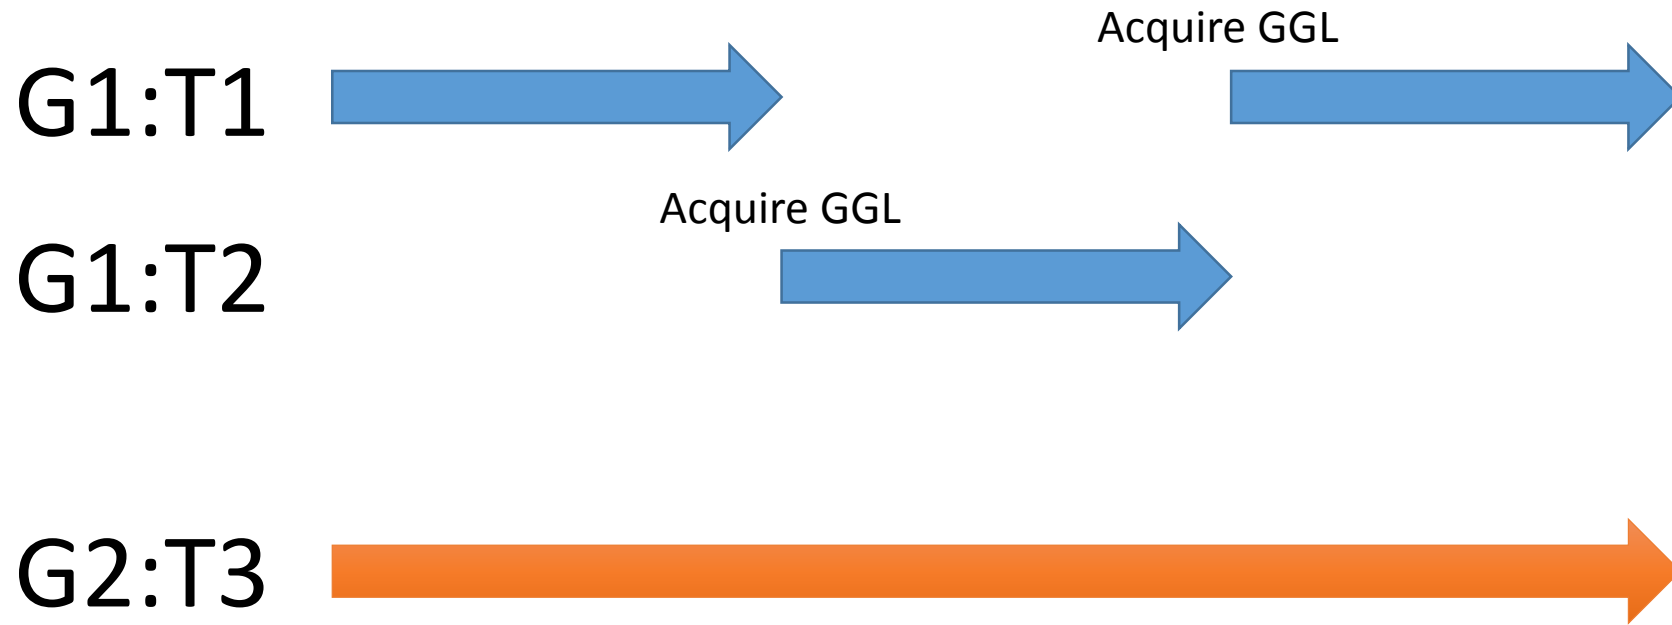
Guild: New concurrency abstraction

- Guild has at least one thread (and a thread has at least one fiber)
 - Threads in different guilds can run in parallel
 - Threads in a same guild can not run in parallel because of GVL (or GGL: Giant Guild Lock)

Each Guild has threads



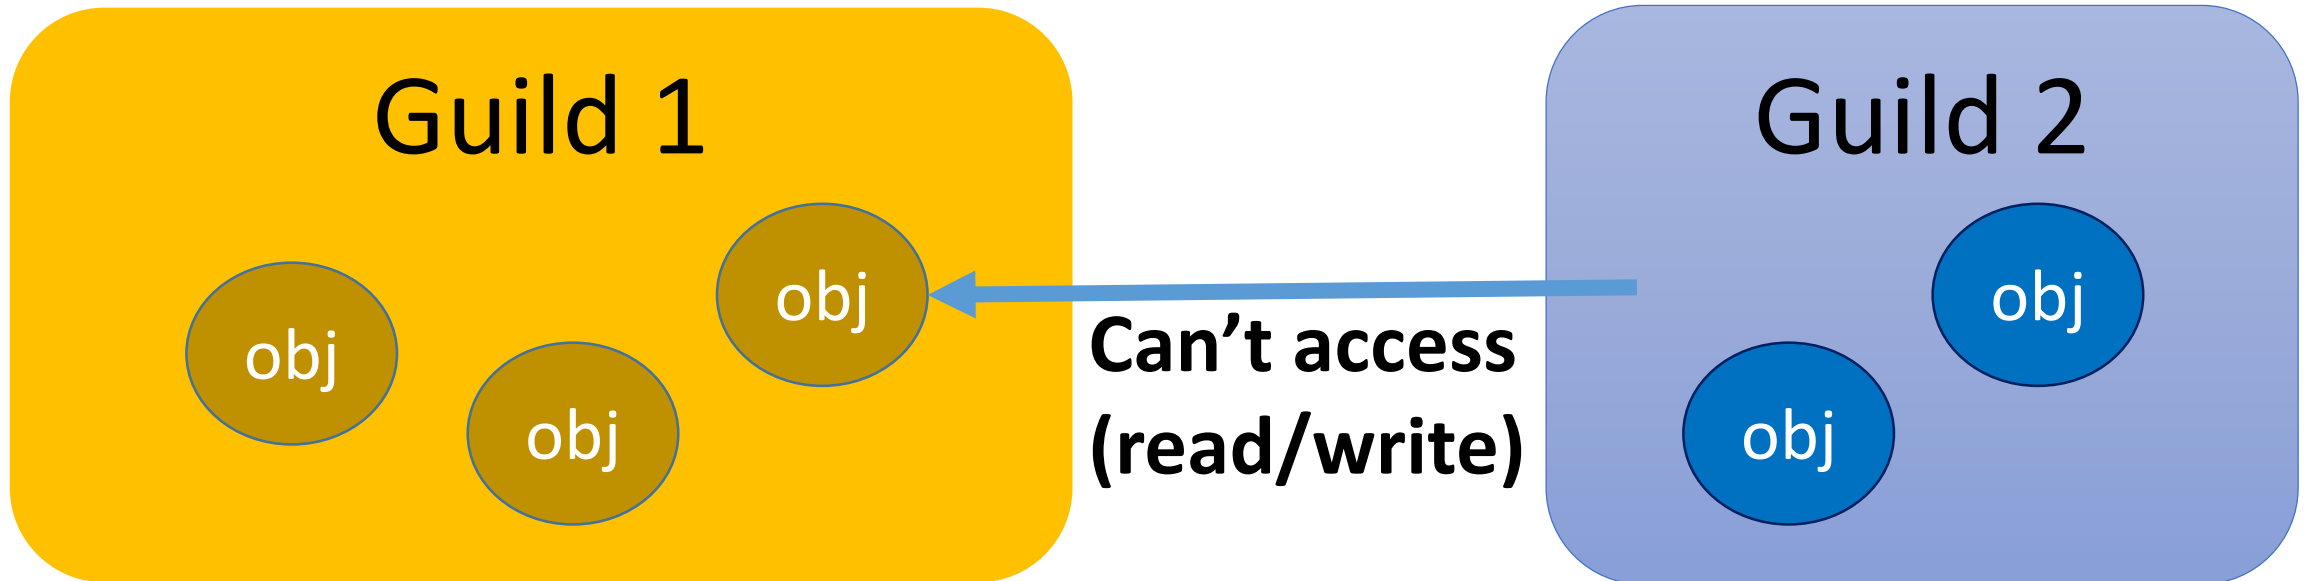
Threads in different guilds can run in Parallel



Guild and objects:

All objects have their own membership

- All of mutable objects should belong to only one Guild (all mutable objects are member of one guild)
- Other guilds can not touch objects



Inter Guilds communication

- We have Guild::Channel to communicate each other
- 2 communication methods
 1. **Copy**
 2. **Transfer membership** or **Move** in short


Copy using Channel

- `Guild::Channel#transfer(obj)` send **deep copied** object(s) to a destination guild.
- dRuby and multi-process system use this kind of communication

Copy using Channel

channel.transfer(o1)

o1 = channel.receive



Move using Channel


[New technique!!]

- `Guild::Channel#transfer_membership(obj)` change the membership of object(s)
 - Leave from the source guild
 - Join to the destination guild
- Accessing from the source guild is invalidated
 - Cause exceptions and so on
 - ex) `obj = "foo"`
`ch.transfer_membership(obj)`
`obj.upcase` `#=> Error!!`
`p(obj)` `#=> Error!!`

Move using Channel

`channel.transfer_membership(o1)`

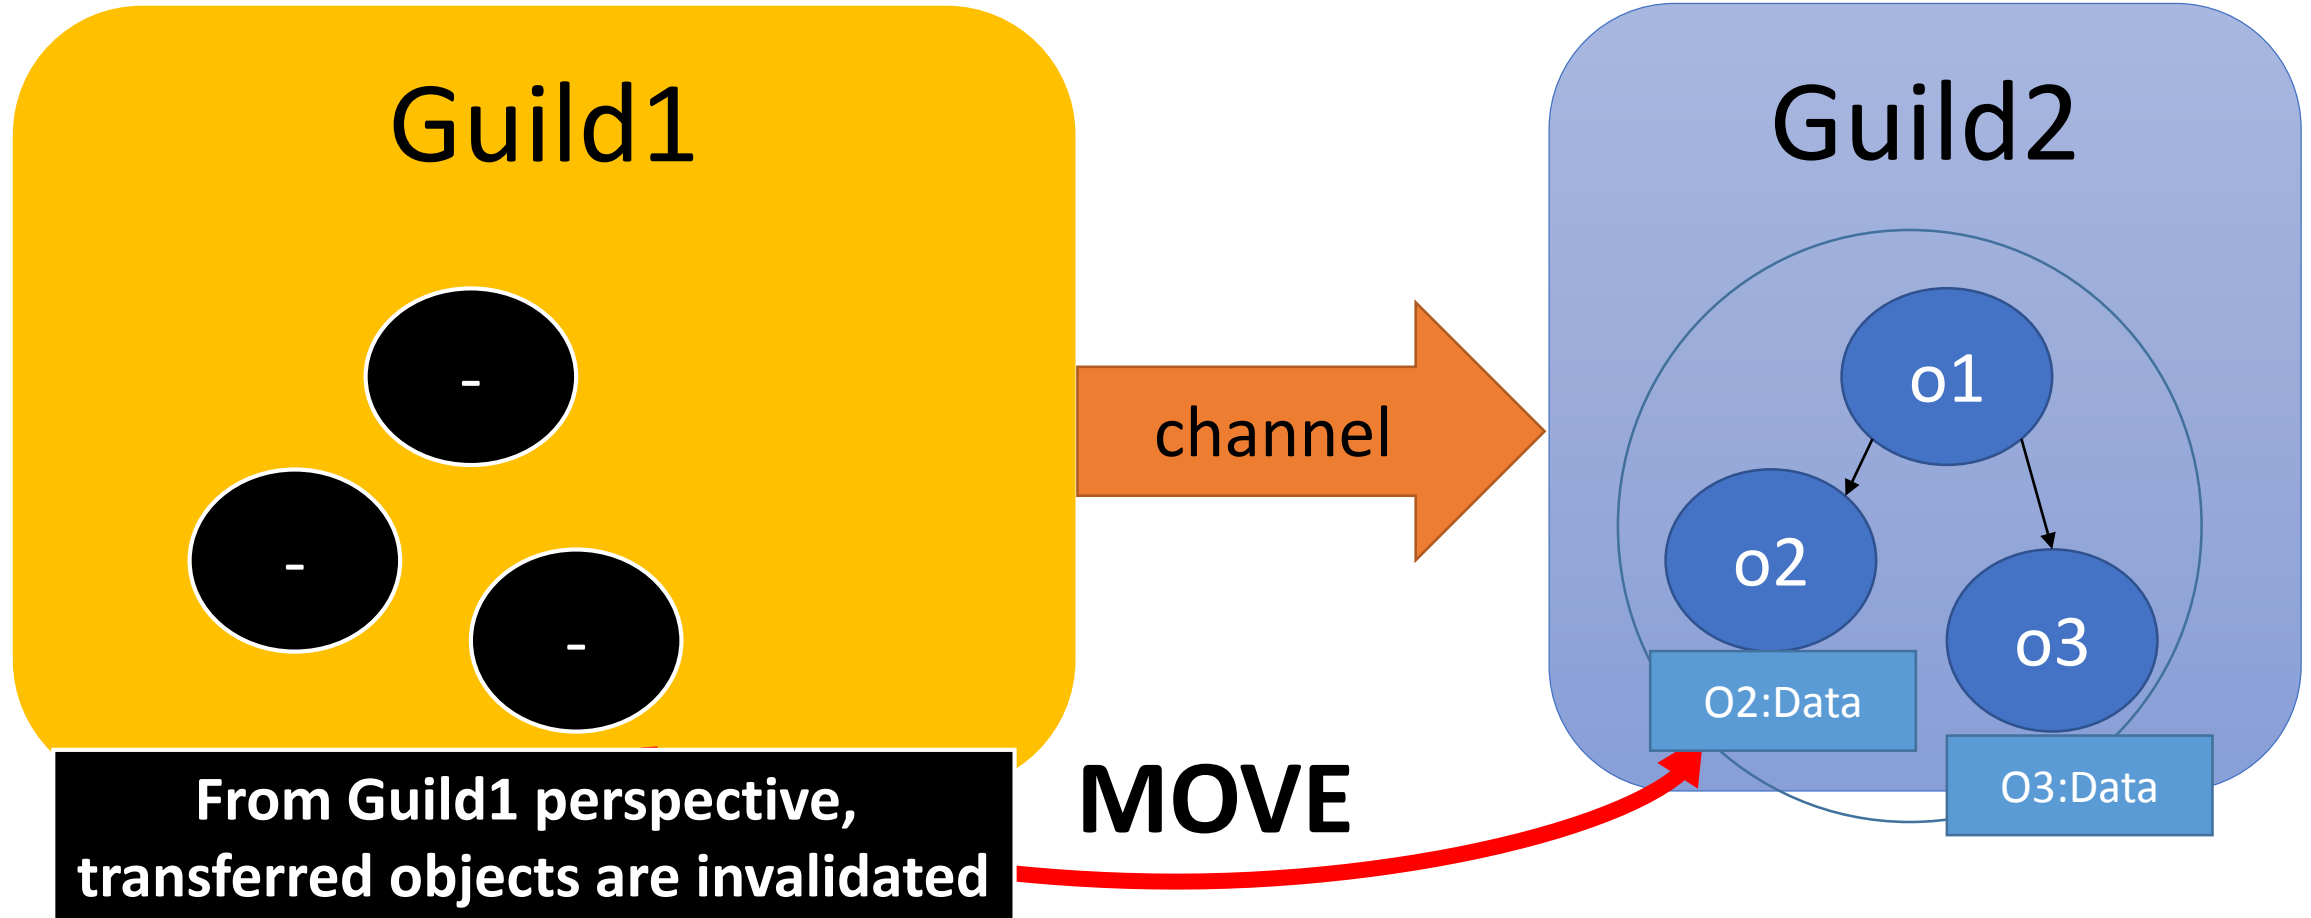
`o1 = channel.receive`



Move using Channel

`channel.transfer_membership(o1)`

`o1 = channel.receive`



Sharing immutable objects

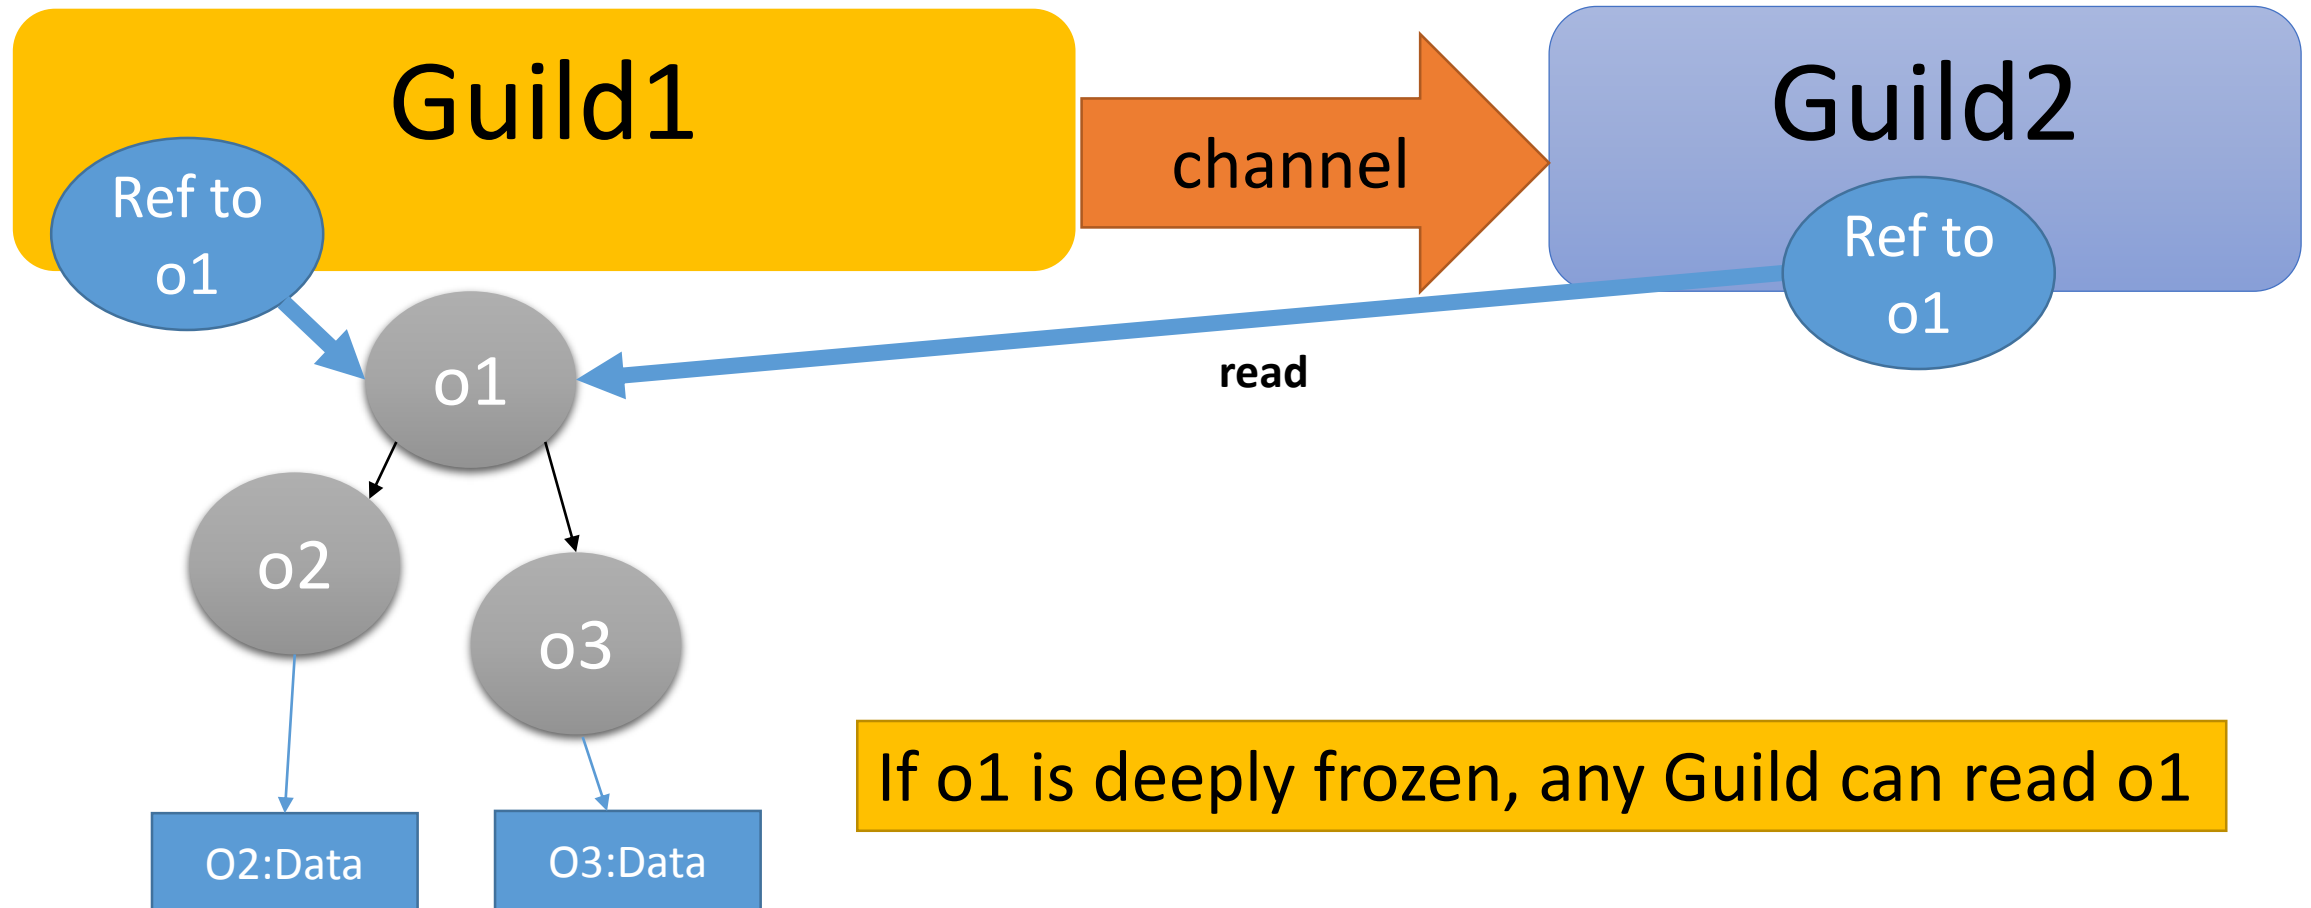
- Deeply frozen objects can be shared with any guilds
 - `[1, 2, 3].freeze`: **deeply frozen**
 - `[1, Object.new, 3].freeze`: not deeply frozen because 2nd element is not frozen
- We only need to send references (very lightweight, like thread-programming)
- **Numeric objects, symbols, true, false, nil** are immutable and deeply frozen (from Ruby 2.0, 2.1, 2.2).

Share immutable objects

We can share reference to immutable objects

channel.transfer(o1)

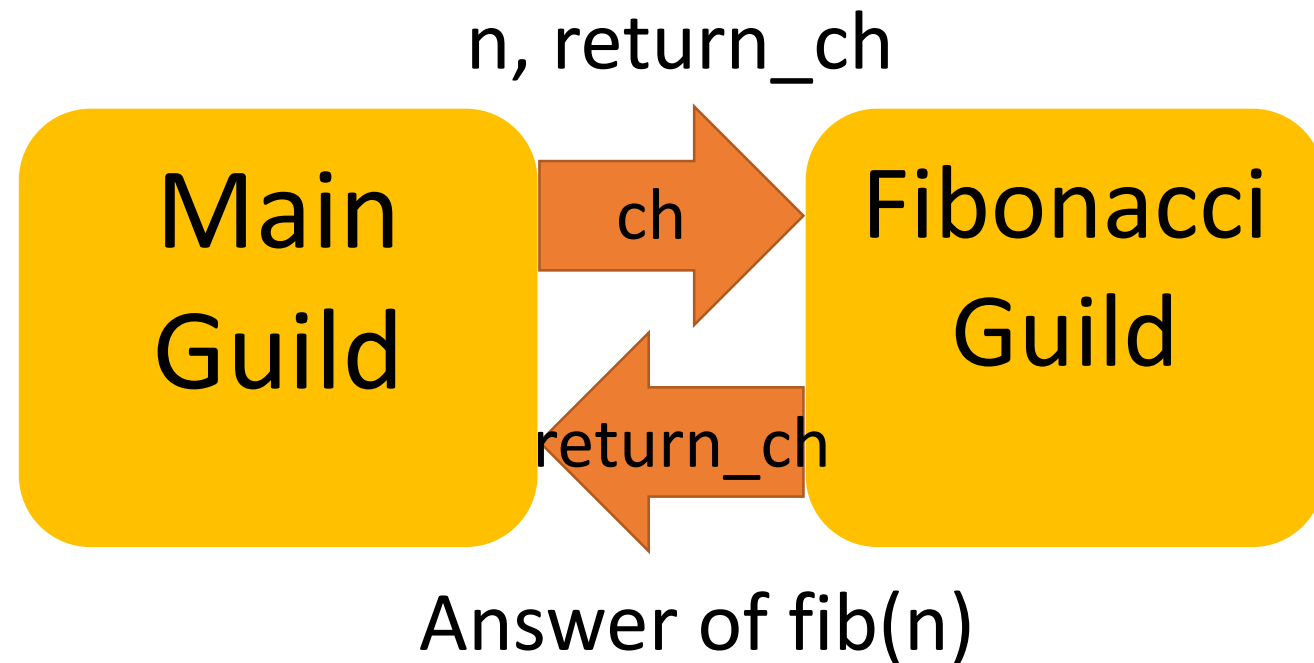
o1 = channel.receive



Use-case 1: master – worker type

```
def fib(n) ... end
g_fib = Guild.new(script: %q{
  ch = Guild.default_channel
  while n, return_ch = ch.receive
    return_ch.transfer fib(n)
  end
})
```

```
ch = Guild::Channel.new
g_fib.transfer([3, ch])
p ch.receive
```



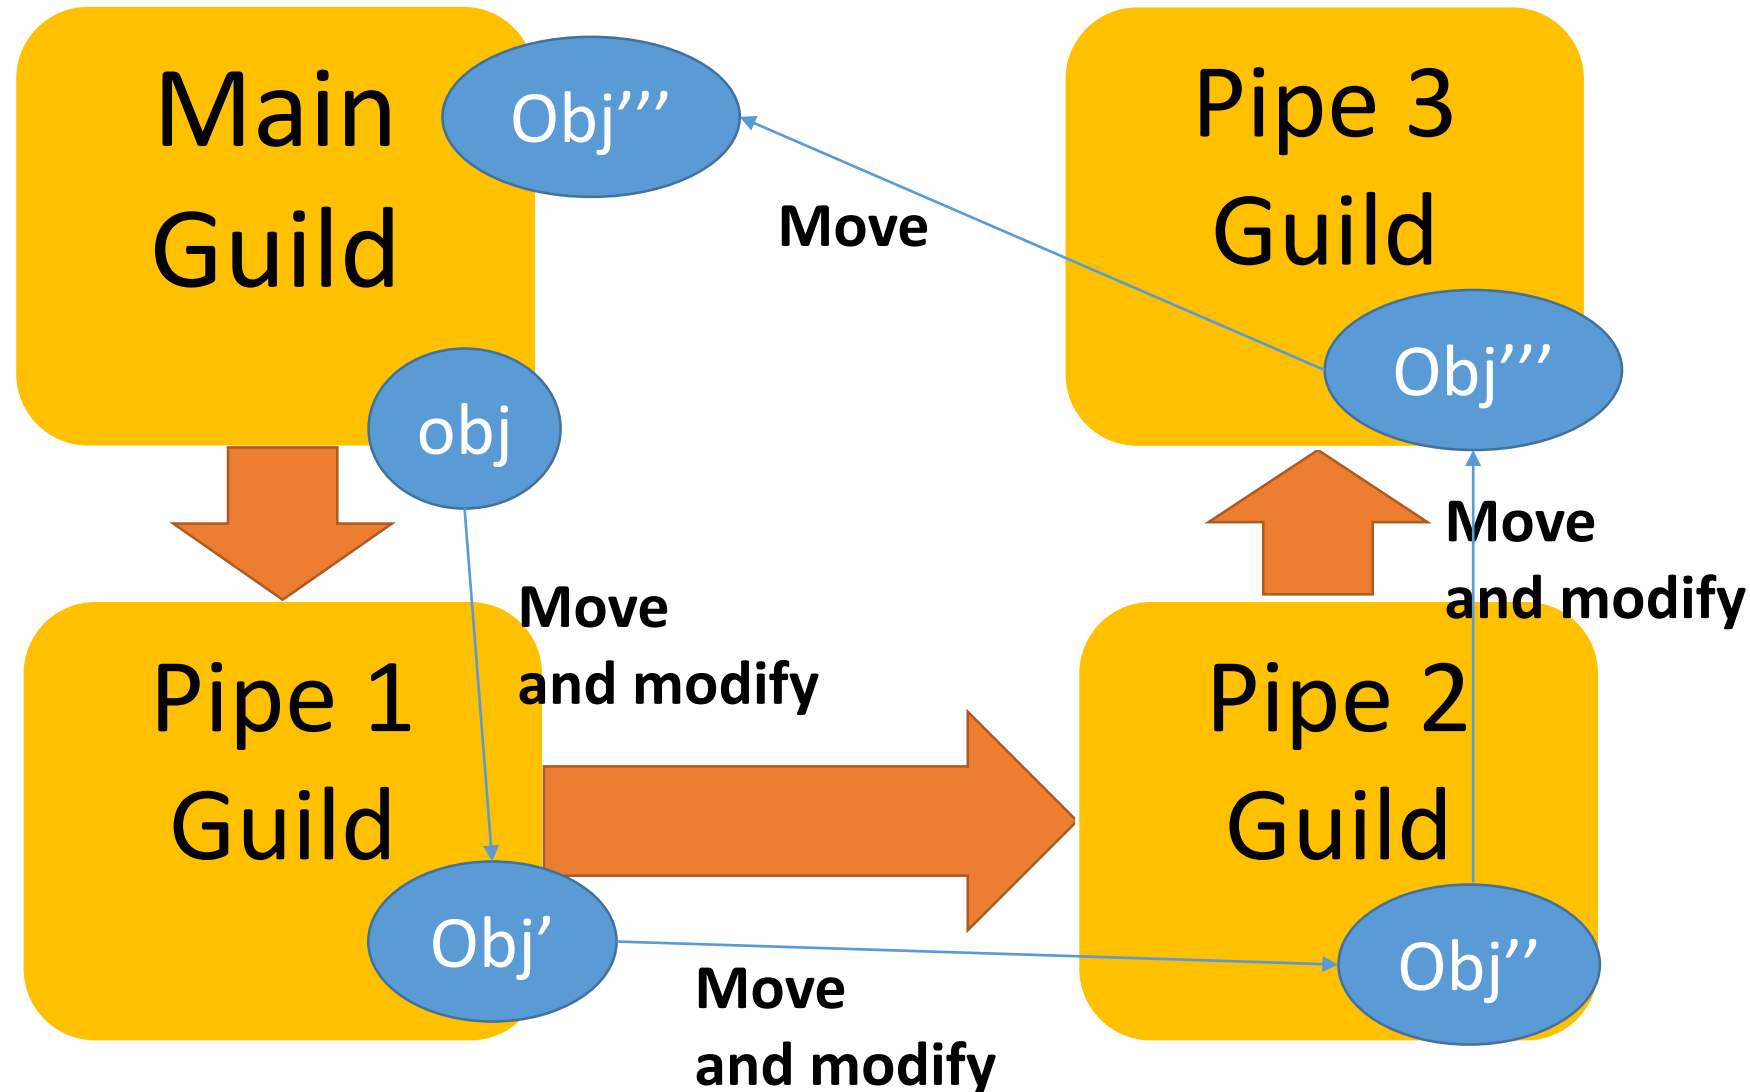
NOTE: Making other Fibonacci guilds, you can compute fib(n) in parallel

Use-case 2: pipeline

```
result_ch = Guild::Channel.new
g_pipe3 = Guild.new(script: %q{
  while obj = Guild.default_channel.receive
    obj = modify_obj3(obj)
  }
  Guild.argv[0].transfer_membership(obj)
end
}, argv: [result_channel])
g_pipe2 = Guild.new(script: %q{
  while obj = Guild.default_channel.receive
    obj = modify_obj2(obj)
  }
  Guild.argv[0].transfer_membership(obj)
end
}, argv: [g_pipe3])
g_pipe1 = Guild.new(script: %q{
  while obj = Guild.default_channel.receive
    obj = modify_obj1(obj)
  }
  Guild.argv[0].transfer_membership(obj)
end
}, argv: [g_pipe2])

obj = SomeClass.new


g_pipe1.transfer_membership(obj)
obj = result_ch.receive
```



Use-case: Bank example

```
g_bank = Guild.new(script: %q{
  while account_from, account_to, amount,
    ch = Guild.default_channel.receive
    if (Bank[account_from].balance < amount)
      ch.transfer :NOPE
    else
      Bank[account_to].balance += amount
      Bank[account_from].balance -= amount
      ch.transfer :YEP
    end
  end
})
...
```


Only bank guild maintains bank data



Use-case:

Introduce special data structure

- Ideas of special data structure to share mutable objects
 - Use external RDB
 - In process/external Key/value store
 - Software transactional memory
 - ...



Summary of use cases

- Making multiple workers and compute in parallel
 - Requests and responses are communicate via channels
 - You can send it with copy or move
 - Maybe web application can employ this model
- Making Pipeline structures and compute in parallel
 - Each task has own Guild
 - Receive target object, modify it and send it next pipeline
 - You will send it with move (transfer membership)
 - It will help applications like applying several filters for input data
- Own responsibility by one Guild
 - All accesses are managed by one responsible Guild
 - If you want to share mutable objects, we need special data structures
 - External RDBs or key/value stores are also good idea for this purpose

Communication strategy

[Upper is better]

- Passing deeply frozen objects
- Copy mutable objects
- If you have performance problem, move (transfer membership) mutable objects
- If you have performance problem too, use special data structure to share mutable objects

Compare between Thread model and Guild model

- On threads, it is **difficult to find out** which objects are shared mutable objects
- On Guilds, there are **no shared mutable objects**
 - If there are special data structure to share mutable objects, we only need to check around this code
 - Encourage “Safe” and “Easy” programming

Compare between Thread model and Guild model

- On threads, inter threads communication is very fast.
- On Guilds, inter guilds communication introduce overhead
 - “Move” (transfer membership) technique can reduce this kind of overheads

Trade-off: Performance v.s. Safety/Easily
Which do you want to choose?

Digression: The name of “Guild”

- “Guild” is good metaphor for “object’s membership”
- Check duplication
 - First letter is not same as other similar abstractions
 - For variable names
 - P is for Processes, T is for Threads, F is for Fibers
 - There are no duplicating top-level classes and modules in all of rubygems

Implementation of “Guild”

- How to implement inter Guilds communication
- How to isolate process global data

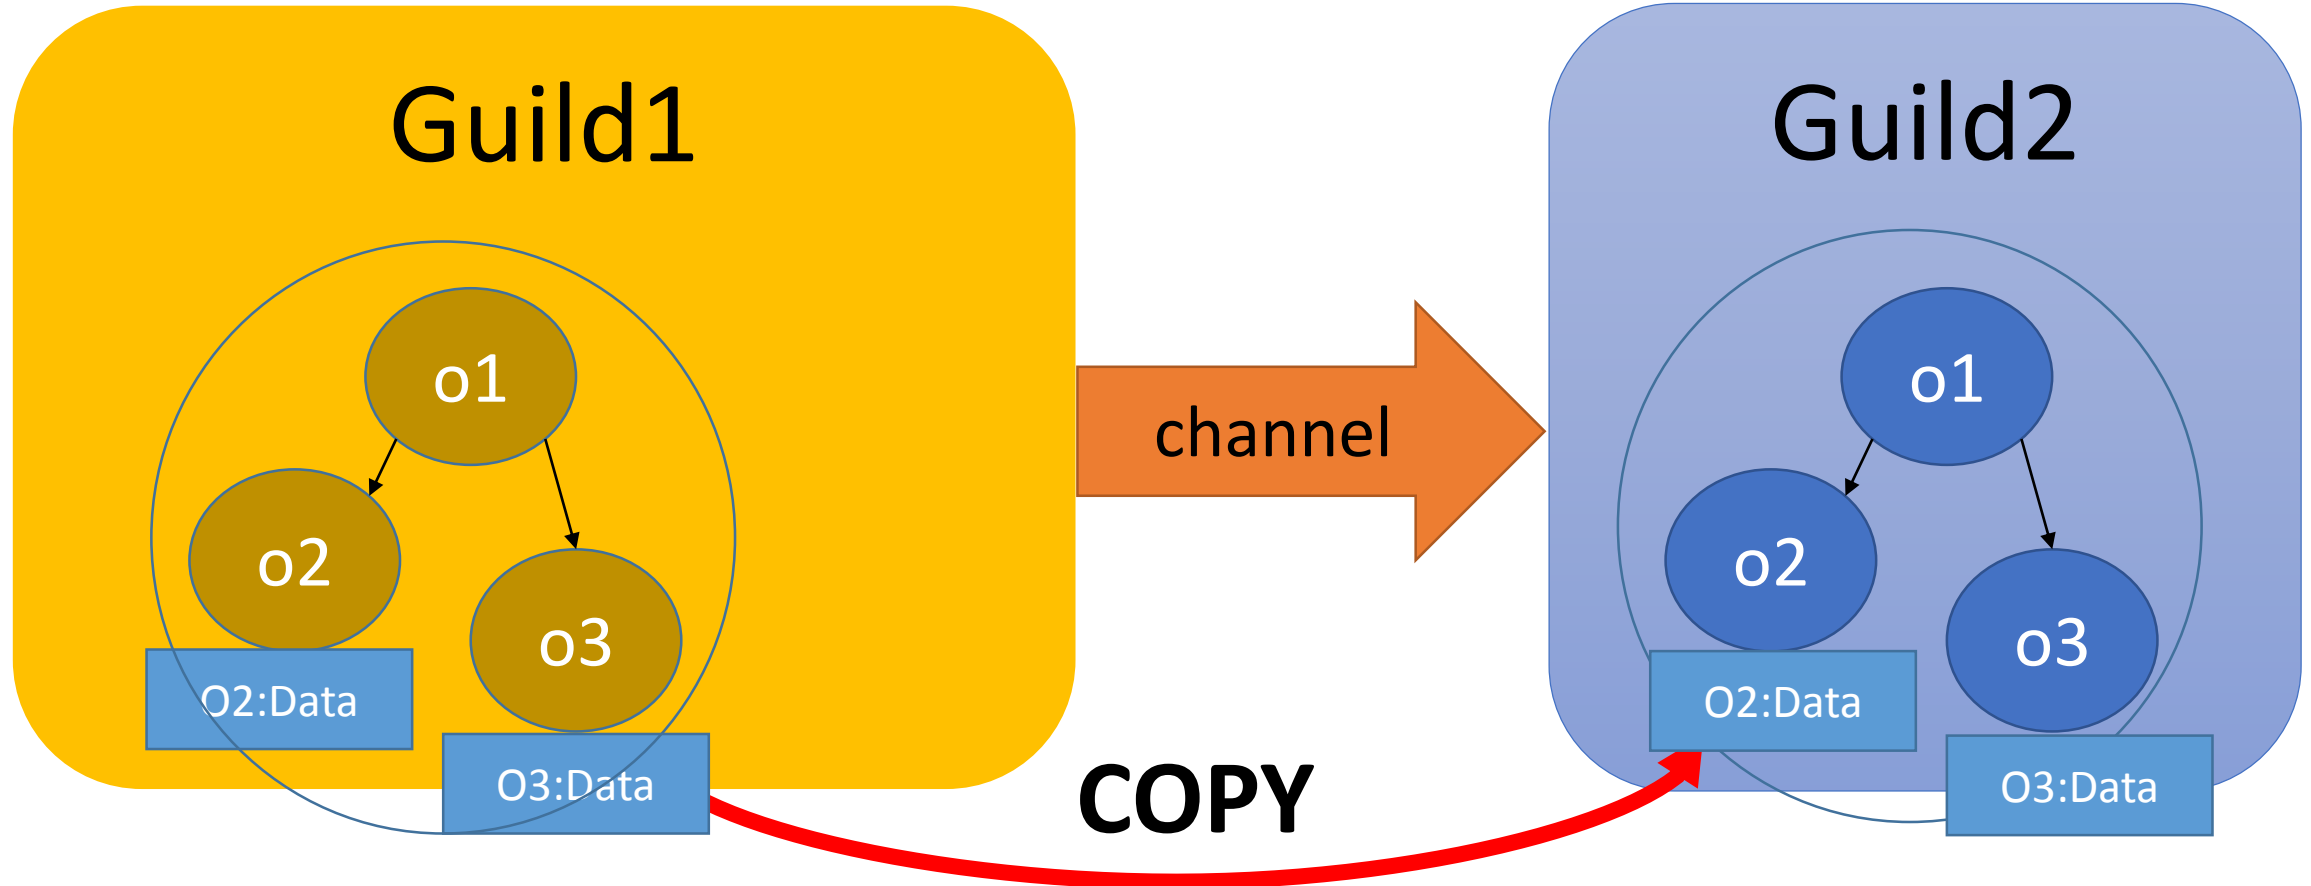
How to implement inter Guilds communication

- Copy
- Move (transfer membership)

Copy using Channel

channel.transfer(o1)


o1 = channel.receive



Copy using Channel Implementation

`channel.transfer(o1)`

`o1 = channel.receive`




Copy using Channel Implementation

`channel.transfer(o1)`

We can use CoW
technique for data


`o1 = channel.receive`



Move using Channel

`channel.transfer_membership(o1)`

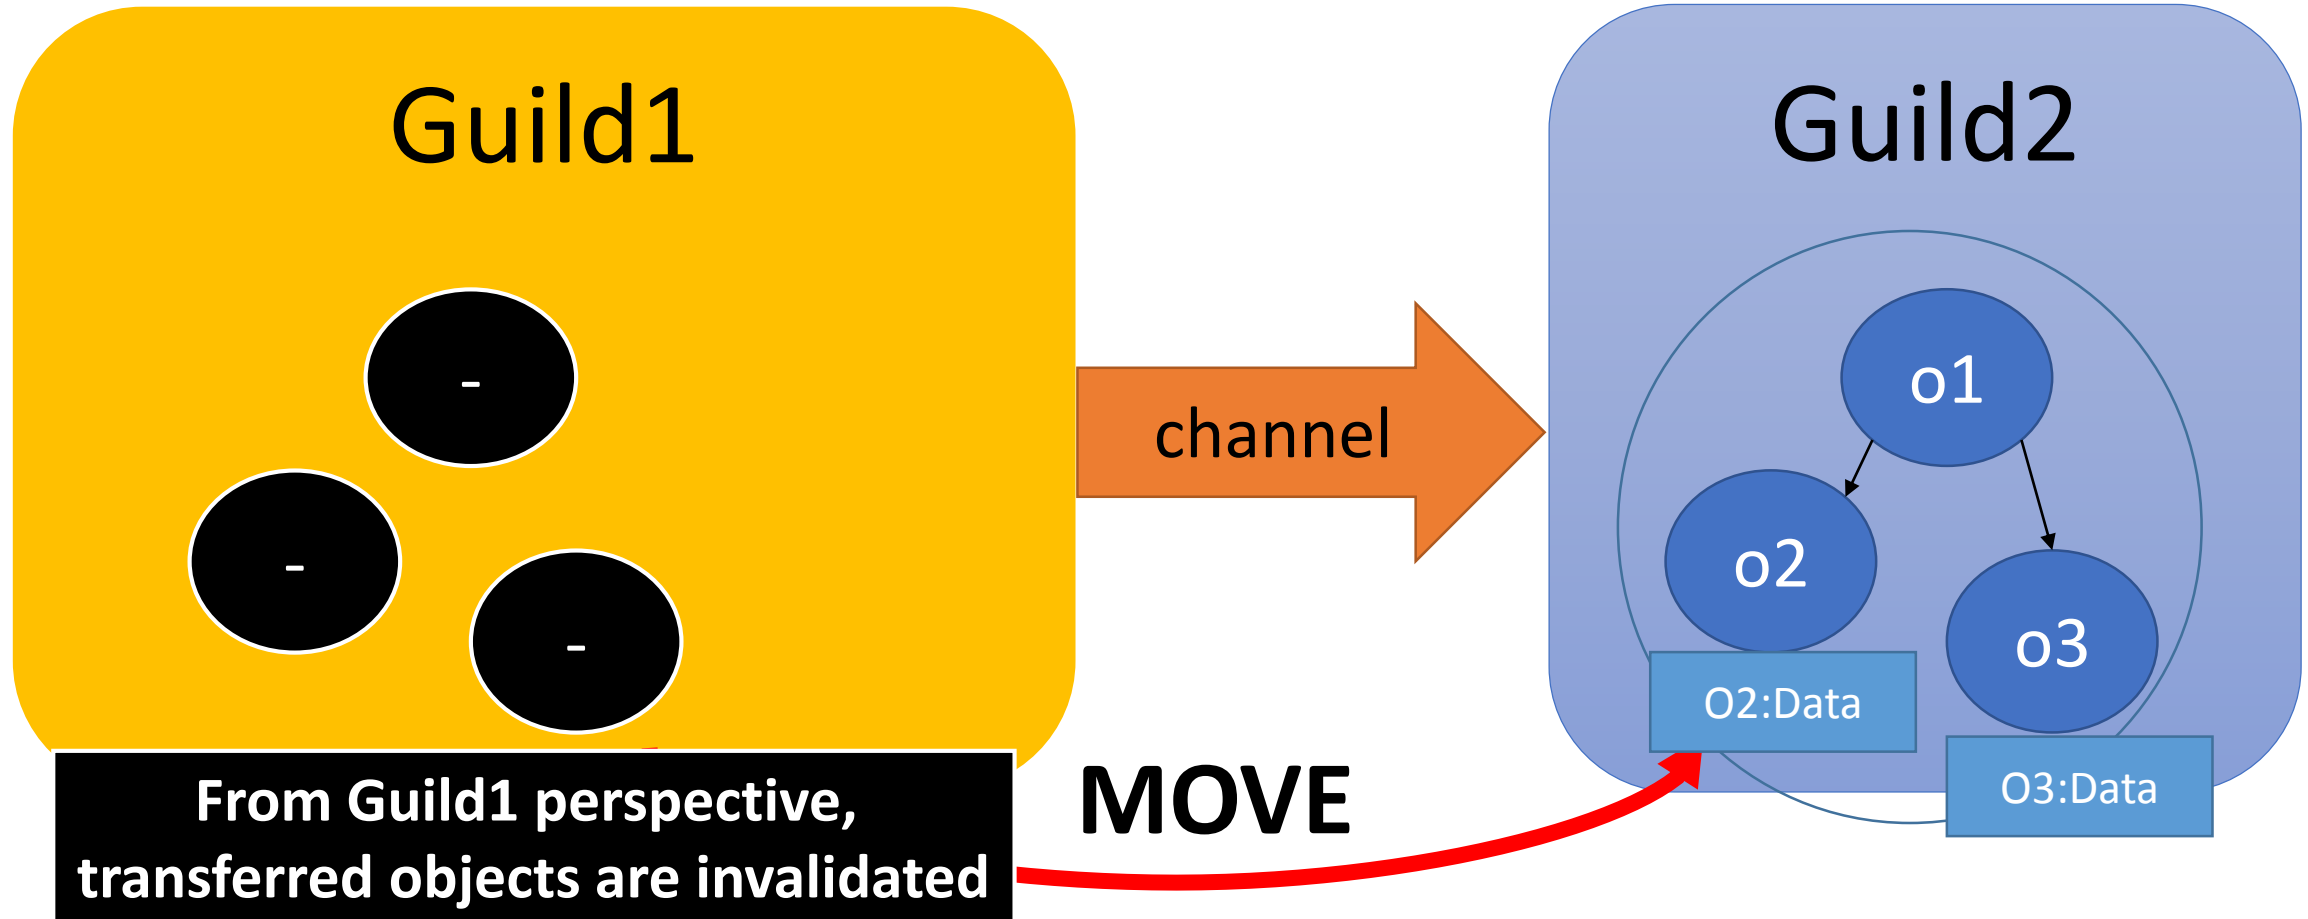
`o1 = channel.receive`



Move using Channel

`channel.transfer_membership(o1)`

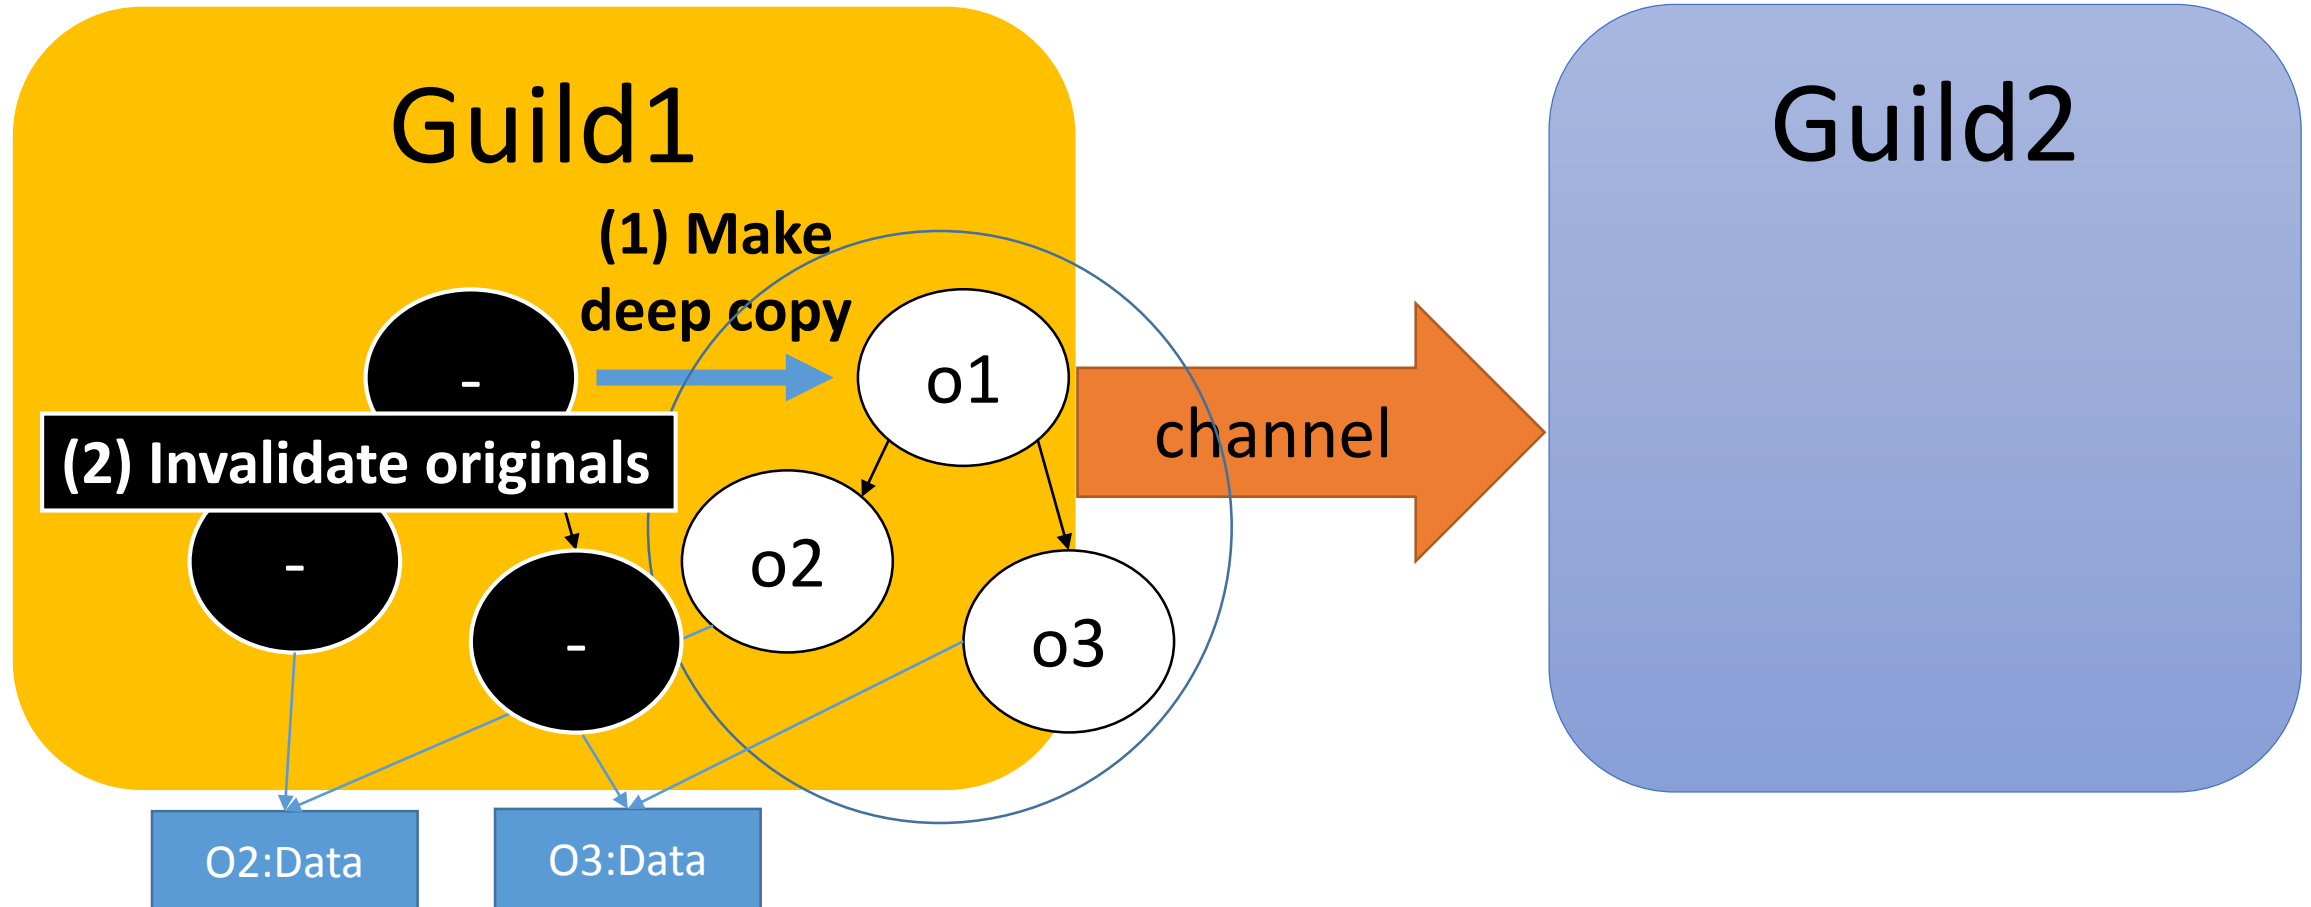
`o1 = channel.receive`



Move using Channel Implementation

`channel.transfer_membership(o1)`


`o1 = channel.receive`



Move using Channel Implementation

`channel.transfer_membership(o1)`

`o1 = channel.receive`



Ruby global data

- Global variables (\$foo)
 - Change them to Guild local variables
- Class and module objects
 - Share between guilds
- Class variables
 - Change them to guild local. So that it is guild/class local variables
- Constants
 - Share between guilds
 - However if assigned object is not a deeply frozen object, this constant is accessed only by setting guilds. If other guilds try to access it, then cause error.
- Instance variables of class and module objects
 - Difficult. There are several approaches.
- Proc/Binding objects
 - Make it copy-able with env objects or env independent objects
- ObjectSpace.each_object
 - OMG

Keep compatibility with Ruby 2

Interpreter process global data

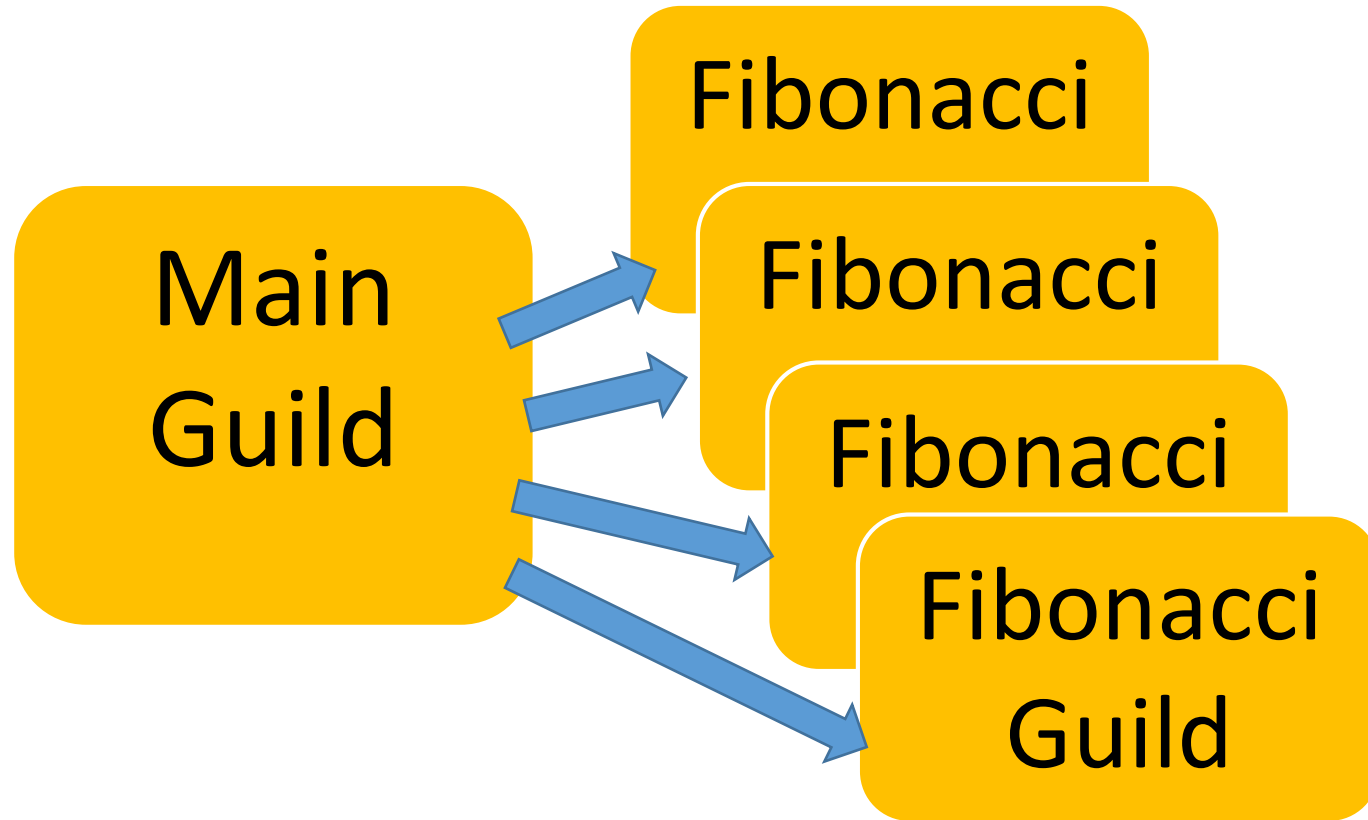
- GC/Heap
 - Share it. Do stop the world parallel marking- and lazy concurrent sweeping.
 - Synchronize only at page acquire timing. No any synchronization at creation time.
- Inline method cache
 - To fill new entry, create an inline cache object and update atomically.
- Tables (such as method tables and constant tables)
 - Introduce mutual exclusions.
- Current working directory (cwd)
 - Each guild should have own cwd (using openat and so on).
- Signal
 - Design new signal delivery protocol and mechanism
- C level global variables
 - Avoid them.
 - Main guild can use C extensions depends on them
- Current thread
 - Use TLS (temporary), but we will change all of C APIs to receive context data as first parameter in the future.

Performance evaluation

- On 2 core virtual machine
 - Linux on VirtualBox on Windows 7
- Now, we can't run Ruby program on other than main guild, so other guilds are implemented by C code

Performance evaluation

Simple numeric task in parallel




	Execution time (sec)
Single-Guild	19.45
Multi-Guild	10.45

Total 50 requests to compute fib(40)
Send 40 (integer) in each request

Performance evaluation

Copy/Move




Total 100 requests to compute sum of array
Send (1..10_000_000).to_a in each request

	Execution time (sec)
Single-Guild	1.00
Multi/ref	0.64
Multi/move	4.29
Multi/copy	5.16

Too slow!!
Because "move" need to
check all of elements

Performance evaluation

Copy/Move



	Execution time (sec)
Single-Guild	1.00
Multi/ref	0.64
Multi/move	0.64

**If we know this array only has deeply frozen objects,
we don't need to check all elements => special data structure**

Check our goal for Ruby 3


- **We need to keep compatibility** with Ruby 2.
 - Only in main guild, it is compatible.
- We shouldn't consider about locks any more.
 - Only using copy and move, we don't need to care locks
- We can share objects with copy, but copy operation should be fast.
 - Move (transfer membership) idea can reduce overhead
- We should share objects if we can.
 - We can share deeply frozen objects easily
- We can provide special objects to share mutable objects like Clojure if we really need speed.
 - Yes, we can provide.

Summary

- Introduce “why threads are very difficult”
- Propose new concurrency abstraction “Guild” for Ruby 3
 - Not implemented yet, but I show key ideas

Thank you for your attention

Koichi Sasada
<ko1@heroku.com>



Approach comparison

	Process/MVM	Place (Racket)	Guild (copy/move)	Thread
Heap	Separate	Separate	Share	Share
Communication Mutable objects	Copy	Copy	Copy/Move	Share
Communication Frozen object	Copy	Share (maybe)	Share	Share
Lock	Don't need	Don't need	(mostly) Don't need	Required
ISeq	Copy	Share	Share	Share
Class/Module	Copy	Copy (fork)	Share	Share